


SAC + β -VAE — Phase 2: Off-Policy Q-Learning

Encoder frozen throughout — no gradient through visual backbone.

No temporal model: each frame processed independently.



$$Q \leftarrow r + \gamma (\min Q' - \alpha \log \pi) \quad \mathcal{L}_\pi = \mathbb{E}[\alpha \log \pi - \min Q]$$